

VISHNU KANTH SURYANARAYAN

AI Engineer · LLM & Multi-Agent Systems · RAG Pipelines

London, UK · +44 7344 701151 · vishnuks0529@gmail.com · Right to work: UK Student Visa, Graduate Visa eligible Oct 2026

github.com/vishnu0529 · linkedin.com/in/vishnu-kanth-suryanarayan-a68851167 · vishnu0529.github.io

PROFESSIONAL SUMMARY

AI Engineer with production experience building multi-agent LLM systems and retrieval-augmented generation (RAG) pipelines using LangChain and LangGraph, alongside a year as a Senior Consultant at Deloitte Digital delivering enterprise digital platforms under SLA, data-governance, and audit requirements. Skilled across the full ML/NLP stack, including PyTorch, HuggingFace Transformers, and Scikit-learn, with MSc research in adversarial robustness testing and model evaluation. Completing an MSc in Artificial Intelligence and Robotics at the University of Hertfordshire (Sep 2026).

CORE SKILLS

Agentic AI & RAG LangChain, LangGraph, Multi-Agent Orchestration, Retrieval-Augmented Generation, Prompt Engineering, RAGAS-style Evaluation, Qdrant, FAISS	Machine Learning & NLP PyTorch, HuggingFace Transformers, Scikit-learn, DistilBERT/BERT Fine-Tuning, Adversarial Robustness (TextFooler, PWWS), Model Evaluation
Backend & APIs Python, FastAPI, Pydantic, PostgreSQL, SQLAlchemy, REST API Design, Microservices	Programming & Data Python, SQL, C#, NumPy, Pandas, Statistical Analysis
MLOps & DevOps Docker, Git, GitHub Actions (CI/CD), pytest, Render, Linux	AI-Native Tooling Google Gemini API, Anthropic Claude API, GitHub Copilot, Cursor

PROFESSIONAL EXPERIENCE

Senior Consultant, Enterprise Digital Engineering · Deloitte Digital, Bangalore, India

Jun 2024 – Sep 2025

- Owned end-to-end delivery of enterprise-scale digital platforms across CMS and UI systems for global clients, building production system ownership, SLA adherence, and large-scale deployment discipline directly applicable to AI product delivery.
- Built and maintained data processing and content workflow automations, applying Python scripting and structured data handling techniques to reduce manual operational overhead across enterprise content pipelines.
- Led cross-functional teams of designers and engineers shipping high-quality digital assets on tight enterprise deadlines, building stakeholder management and technical delivery skills transferable to AI engineering teams.
- Championed adoption of AI-assisted development tooling (GitHub Copilot, LLM-assisted workflows) within the team, accelerating delivery cycles and improving code review quality.
- Operated within compliance-driven enterprise environments with data governance, access controls, and audit requirements, foundational experience for deploying AI systems in regulated industries.

Front-End Engineer / UI-UX Designer · PinkLemonade Communications, Bangalore, India

Nov 2021 – Jul 2023

- Designed and delivered user interfaces for SaaS and e-commerce products across 6+ client projects, owning end-to-end delivery from discovery through QA in agile sprint cycles.
- Collaborated with backend engineers to ensure design-implementation fidelity, developing structured problem-solving and cross-functional communication skills foundational to AI product work.

Independent IT Consultant, Sole Trader · Bangalore, India

2018 – 2021

- Provided general IT and software consulting services to small-business clients as a self-employed contractor.

KEY PROJECTS

Enterprise Knowledge Assistant: Two-Agent LangGraph RAG Platform [2026]

Python, LangGraph, Qdrant, FastAPI, Docker · github.com/vishnu0529/enterprise-rag-assistant

- Built for enterprise knowledge-management use cases where hallucinated answers carry real business risk, grounding every answer in cited source documents rather than the model's own prior knowledge.
- Designed a two-agent system: a Retrieval Strategist that plans searches, including multi-hop query decomposition for comparison-style questions, and a Drafting Agent that writes answers strictly from retrieved evidence, coordinated through LangGraph shared state.
- Added a critique node running a corrective self-RAG loop that re-plans ungrounded answers instead of shipping them, with a RAGAS-style evaluation suite (faithfulness, answer relevancy, context precision/recall).
- Shipped with Docker, a GitHub Actions CI pipeline, API-key authentication, and 30 passing tests; dual LLM-provider abstraction (Google Gemini, Anthropic Claude).

AI Job Finder Bot: Multi-Agent LangGraph Job Search Assistant [2026]

Python, LangGraph, Google Gemini API, FAISS, Streamlit · github.com/vishnu0529/ai-job-finder-bot

- Built for the UK recruitment domain: a LangGraph state-machine agent that scores real job postings and checks live UK sponsor-register status, addressing a real compliance concern for visa-dependent job seekers.
- Drafts a tailored cover letter and critiques its own draft, looping back to revise when the critic scores it below threshold; FAISS-backed cross-session memory avoids repeating prior cover-letter openers.

AI Resume Matcher: Production Agentic LLM System [2026]

Google Gemini, Anthropic Claude, FastAPI, Pydantic, Render · github.com/vishnu0529/ai-resume-matcher

- Built for the HR-tech and talent-acquisition domain: a 4-step agentic reasoning pipeline (skill extraction, gap analysis, content generation, strategy synthesis) that analyses a CV against any job description.
- Production FastAPI backend with schema-enforced (Pydantic) API contracts and dual LLM-provider support, deployed live with fully automated CI/CD and 23 passing tests.

Phishing Detection: Classical ML vs Transformer NLP [MSc Project, 2026]

Python, PyTorch, HuggingFace Transformers, DistilBERT, Scikit-learn · [University of Hertfordshire](https://www.hertfordshire.ac.uk/research/centres-and-groups/centre-for-artificial-intelligence)

- Cybersecurity-focused NLP research: benchmarked Naive Bayes, Random Forest, and a fine-tuned DistilBERT transformer on phishing-email text classification, reaching 99.4% accuracy / 0.994 F1 with DistilBERT vs 97.9% for the strongest classical baseline.
- Extended evaluation to adversarial robustness (TextFooler, PWWS attacks): the transformer's accuracy gain did not translate into a robustness gain, an evidence-based model-selection finding relevant to deploying NLP classifiers in adversarial, real-world security settings.

EDUCATION

MSc Artificial Intelligence and Robotics · University of Hertfordshire · 2025 – Sep 2026

BEng Electronics and Communication Engineering · East West Institute of Technology · 2014 – 2018

CERTIFICATIONS

Certificate of Attendance, From Zero to Hero: Federated AI in Healthcare Systems · St John's College, Cambridge · Aug 2026

ParsAI Workshop Certificate, Prompt Engineering · University of Hertfordshire (Dr. Abolfazl Zaraki) · Apr 2026