

VISHNU KANTH SURYANARAYAN

AI Engineer & Independent AI Consultant · Agentic AI / LangGraph · RAG Systems

vishnuks0529@gmail.com | +44 7344 701151 | London, UK | Right to Work: Student Visa · Graduate Visa eligible Oct 2026
github.com/vishnu0529 | linkedin.com/in/vishnu-kanth-suryanarayan-a68851167 | vishnu0529.github.io

PROFESSIONAL SUMMARY

Independent AI Consultant and AI Engineer completing an MSc in Artificial Intelligence & Robotics at the University of Hertfordshire (Sep 2026). Currently designs and ships production multi-agent LLM systems end-to-end: a two-agent LangGraph RAG platform (a Retrieval Strategist that plans multi-hop searches and a Drafting Agent that writes from evidence only, with a corrective/self-RAG critique loop and Qdrant-backed cross-session memory), and a LangGraph job-search agent with Gemini-based scoring, self-critiquing cover-letter generation, FAISS-backed memory, and real UK sponsor-register checks — both deployed live with full CI/CD. MSc research separately compared classical (Naive Bayes, Random Forest) and transformer-based (DistilBERT) models for text classification, including adversarial robustness testing (TextFooler, PWWS). Also built a federated learning pipeline (Flower) predicting hospital readmissions across simulated hospital sites without centralizing patient data, and completed a two-day Cambridge workshop on federated AI in healthcare systems. A year at Deloitte Digital, delivering enterprise platforms under SLA, data-governance, and audit requirements, gave direct experience with the compliance-driven, stakeholder-facing delivery regulated AI products demand. Seeking an AI Engineering role with full-time availability from October 2026.

CORE COMPETENCIES

Agentic AI & RAG	LangChain · LangGraph · Multi-Agent Orchestration · Stateful Checkpointing · Corrective / Self-RAG · Qdrant / FAISS · RAGAS-style Evaluation · Prompt Engineering	Machine Learning & NLP	PyTorch · HuggingFace Transformers · Scikit-learn · DistilBERT / BERT Fine-Tuning · Adversarial Robustness (TextFooler, PWWS) · Model Evaluation
Backend & APIs	FastAPI · Pydantic · PostgreSQL / SQLAlchemy · REST API Design · Schema Validation · Microservices	Programming & Data	Python · SQL · C# (OOP) · NumPy · Pandas · EDA · Statistical Analysis
MLOps & DevOps	Git · GitHub · Docker · CI/CD (GitHub Actions) · pytest · Linux · Azure (ASP.NET Core project) · Jupyter	AI-Native Tooling	Google Gemini API · Anthropic Claude API · GitHub Copilot · Cursor

PROFESSIONAL EXPERIENCE

Independent AI Consultant

· Remote / London, UK

Sep 2026 – Present

- Design and ship production multi-agent LLM systems end-to-end: LangGraph state machines with critique/revision loops, cross-session memory (Qdrant, FAISS), dual-LLM-provider abstraction (Gemini, Claude), and API-key-gated FastAPI backends deployed on Render with Streamlit front-ends.
- Own the full lifecycle — architecture, testing (pytest, 50+ tests across projects), CI/CD, live deployment, and post-launch fixes (rate limiting, quota handling, memory-constrained deployment via fastembed/ONNX) — found and resolved through real production testing rather than local-only development.

Senior Consultant, Enterprise Digital Engineering · Deloitte Digital, Bangalore, India Jun 2024 – Sep 2025

- Owned end-to-end delivery of enterprise-scale digital platforms across CMS and UI systems for global clients — developing production system ownership, SLA adherence, and large-scale deployment discipline

directly applicable to AI product delivery.

- Built and maintained data processing and content workflow automations, applying Python scripting and structured data handling techniques to reduce manual operational overhead across enterprise content pipelines.
- Led cross-functional teams of designers and engineers shipping high-quality digital assets on tight enterprise deadlines — building stakeholder management and technical delivery skills transferable to AI engineering teams.
- Championed adoption of AI-assisted development tooling (GitHub Copilot, LLM-assisted workflows) within the team, accelerating delivery cycles and improving code review quality.
- Operated within compliance-driven enterprise environments with data governance, access controls, and audit requirements — foundational experience for deploying AI systems in regulated industries.

Front-End Engineer / UI-UX Designer · PinkLemonade Communications, Bangalore, India *Nov 2021 – Jul 2023*

- Designed and delivered user interfaces for SaaS and e-commerce products across 6+ client projects; owned end-to-end delivery from discovery through QA in agile sprint cycles.
- Collaborated with backend engineers to ensure design–implementation fidelity; developed structured problem-solving and cross-functional communication skills foundational to AI product work.

Independent IT Consultant, Sole Trader

· Bangalore, India

2018 – 2021

- Provided general IT and software consulting services to small-business clients as a self-employed contractor.

KEY PROJECTS

AI Job Finder Bot — Multi-Agent LangGraph Job Search & Application Assistant [2026]

Stack: Python · LangGraph · Google Gemini API · FAISS · Streamlit · SQLite | github.com/vishnu0529/ai-job-finder-bot |

Live: ai-job-finder-bot-bnqepy7obmunbihrkyxv4f.streamlit.app

- Built a LangGraph state-machine agent that scores real job postings, checks UK sponsor-register status, drafts a tailored cover letter, and critiques its own draft — looping back to revise when the critic scores it below threshold.
- Implemented FAISS-backed cross-session memory so the agent recalls past applications and avoids repeating prior cover-letter openers; added shared Gemini rate-limiting/retry and quota-exceeded handling for resilient long-running runs.

Enterprise Knowledge Assistant — Two-Agent LangGraph RAG Platform [2026]

Stack: FastAPI · LangGraph · Qdrant · fastembed · Pydantic · Docker · Render · Streamlit |

github.com/vishnu0529/enterprise-rag-assistant | Live: enterprise-rag-assistant-bpoiuckv3cxhsbkuaay8ve.streamlit.app

- Designed a genuine two-agent split — a Retrieval Strategist that plans searches (including multi-hop query decomposition) and a Drafting Agent that writes answers strictly from retrieved evidence — coordinated through LangGraph shared state.
- Added a critique node implementing a corrective/self-RAG loop: ungrounded answers are routed back to re-plan retrieval rather than returned to the user; cross-session memory persisted via Qdrant.
- Shipped with Docker, a GitHub Actions CI pipeline (lint → test → build), API-key authentication, and 30 passing tests; dual LLM-provider abstraction (Google Gemini / Anthropic Claude).

AI Resume Matcher — Production Agentic LLM System [2026]

Stack: Google Gemini 3.6 Flash · Anthropic Claude · FastAPI · Streamlit · Pydantic · Render · GitHub Actions |

github.com/vishnu0529/ai-resume-matcher | Live: ai-resume-matcher-afsgzlmmlkspynzeebp9w4.streamlit.app

- Shipped a full-stack agentic LLM system with a 4-step reasoning pipeline (skill extraction → gap analysis → content generation → strategy synthesis) and a production FastAPI backend with schema-enforced (Pydantic) API contracts.
- Deployed live on Render with fully automated CI/CD via GitHub Actions and 23 passing tests.

Privacy-Preserving Federated Healthcare Intelligence [2026]

Stack: Python · Flower (flwr) · Scikit-learn · Pandas · Matplotlib | github.com/vishnu0529/Federated_Project

- Built a horizontal federated learning pipeline (Flower) predicting 30-day hospital readmissions across three simulated hospital sites without centralizing patient data — handling non-IID label skew and simulated inter-site network delay, with FedAvg weighted aggregation across training rounds.
- Exported the global model with a standalone inference script and training-loss/accuracy visualisations, directly modelling the constraint clinical platforms face: data that can't leave the source system.

Phishing Detection — Classical ML vs Transformer NLP [2026]

Stack: Python · PyTorch · HuggingFace Transformers · DistilBERT · Scikit-learn · Naive Bayes · Random Forest | MSc Final Project, University of Hertfordshire

- Benchmarked Naive Bayes, Random Forest (TF-IDF), and a fine-tuned DistilBERT transformer on phishing-email text classification, reaching 99.4% accuracy / 0.994 F1 with DistilBERT vs 97.9% / 95.5% for the strongest classical baseline.
- Extended evaluation to adversarial robustness (TextFooler, PWWS attacks): the transformer's accuracy gain did not translate into a robustness gain — an evidence-based model-selection finding, formally written up across a full MSc dissertation.

CERTIFICATIONS

- **Certificate of Attendance — From Zero to Hero: Federated AI in Healthcare Systems**, St John's College, Cambridge (24–25 Aug 2026), hosted by Cambridge Image Analysis (DAMTP, University of Cambridge), PharosAI & Core AI, built on the Flower (flwr) framework. Verify: verified.certifier.com/en/verify/99023372901934

EDUCATION

MSc Artificial Intelligence & Robotics · University of Hertfordshire · 2025 – Sep 2026

Key Modules: Machine Learning · AI Search · Neural Networks · Robotics · Probabilistic AI · AI Systems Design · Robotic Control · 3D Modelling & Simulation

BEng Electronics & Communication Engineering · East West Institute of Technology · 2014 – 2018

ADDITIONAL INFORMATION

Languages	Domain Focus	Availability	Live Systems
English — Full Professional Proficiency	Agentic AI / Multi-Agent Systems · Healthcare AI · NLP / Text Classification · Federated / Privacy-Preserving ML	Part-time (20 hrs/wk) until Sep 2026 · Full-time from Oct 2026	3 live projects — vishnu0529.github.io