

VISHNU KANTH SURYANARAYAN

AI Engineer | Agent Systems, Pilot to Production

vishnuks0529@gmail.com | +44 7344 701151 | London, UK | github.com/vishnu0529 | linkedin.com/in/vishnu-kanth-suryanarayan-a68851167

PROFESSIONAL SUMMARY

AI Engineer: agent systems from pilot to production. I build the layer that decides whether an agent survives a real deployment: evaluation gates, cost per task, data boundaries, human approval on irreversible actions, and a named owner. LangGraph and LangChain orchestration with durable cross-process state; retrieval with enforced citations and a RAGAS-style eval suite; four production-hardening sprints across security, reliability, scalability and LLM cost. Ex-Deloitte Digital Senior Consultant. MSc AI & Robotics, University of Hertfordshire, Sept 2026.

CORE COMPETENCIES

Agent Evaluation & Guardrails	RAGAS Evaluation, Golden Sets & Adversarial Traps, Self-RAG, Grounding & Citation Enforcement, Critique & Re-Planning Loops, Escalation Paths
Agentic AI & LLM Orchestration	LangChain, LangGraph, Multi-Agent Systems, Durable Cross-Process Checkpointing (Postgres), Human-in-the-Loop (interrupt()), OpenTelemetry Tracing, Cross-Agent Memory, RAG, Prompt Engineering, Vector Databases (Qdrant, FAISS)
Production Hardening	Data-Boundary Controls, One-Command Rollback, CI Regression Gates, Cost-per-Task & Latency Tracking, Config/Prompt Versioning, Security & Access Control
Machine Learning & NLP	PyTorch, HuggingFace Transformers, Scikit-learn, DistilBERT/BERT Fine-Tuning, Adversarial Robustness (TextFooler, PWWS), Model Evaluation
Backend & APIs	Python, FastAPI, Pydantic, PostgreSQL, SQLAlchemy, REST API Design, Microservices
MLOps & DevOps	Docker, Git, GitHub Actions (CI/CD), pytest, Render, Linux
AI-Native Tooling	Google Gemini API, Anthropic Claude API, GitHub Copilot, Cursor

PROFESSIONAL EXPERIENCE

Senior Consultant, Enterprise Digital Engineering | Deloitte Digital

Jun 2024 – Sep 2025 | Bangalore, India

- Owned end-to-end delivery of enterprise-scale digital platforms for global clients, meeting SLA, data-governance, and audit requirements.
- Built Python-based data processing and workflow automations, cutting manual work across content pipelines.
- Led cross-functional teams of designers and engineers to ship on tight deadlines.
- Introduced AI-assisted tools (GitHub Copilot, LLM workflows) that sped up delivery and improved code review.
- Worked under strict data governance and audit rules, directly relevant to deploying AI in regulated industries.

Front-End Engineer / UI-UX Designer | PinkLemonade Communications

Nov 2021 – Jul 2023 | Bangalore, India

- Designed interfaces for SaaS and e-commerce products across 6+ client projects, from discovery through QA.

Independent IT Consultant, Sole Trader

2018 – 2021 | Bangalore, India

- Provided IT and software consulting to small-business clients as a self-employed contractor.

KEY PROJECTS

AI Business Automation Hub: Multi-Agent LangGraph Automation Platform (2026)

- Hit and solved a real production failure: process-local agent checkpointing broke across replicas. Rewrote it on Postgres-backed durable checkpointing, keyed per engagement.
- Supervisor agent routes to four specialist agents (Sales, Support, Research, Marketing) over a LangGraph StateGraph, with cross-session memory via a pgvector-backed store.
- Four hardening sprints (security, reliability, scalability, LLM cost); 307 tests passing. Built, hardened and tested, with staging deploy in progress; not yet claimed as in production.

Proposal Response Assistant: Multi-Agent LangGraph RAG for Professional-Services Bid Teams (2026)

github.com/vishnu0529/enterprise-rag-assistant

- Two-agent corrective RAG drafts RFP answers from a firm's own proposals, credentials and rate card, every claim tied to a source.
- Citation enforcement, escalation, and a genuine human-in-the-loop approval gate (LangGraph interrupt()) on commercially sensitive answers.
- Durable Postgres checkpointing verified with a real two-process, SIGKILL demonstration. OpenTelemetry tracing per node. 50-item golden set with 12 deliberate refusal traps, wired into a CI eval gate.
- Scored against a public 15-point Agent Production Readiness Scorecard (self-assessed: 9 demonstrated, 5 partial, 1 open).

AI Job Finder Bot: Multi-Agent LangGraph Job Search Assistant (2026)

github.com/vishnu0529/ai-job-finder-bot

- Scores real UK job postings and checks live UK sponsor-register status. Writes and self-critiques a cover letter, remembering past letters across sessions.

Phishing Detection: Classical ML vs Transformer NLP, MSc Project (2026)

- Benchmarked classical ML against a fine-tuned transformer: DistilBERT reached 99.4% accuracy vs 97.9% for the best classical model. Tested both under adversarial attack (TextFooler, PWWS).

EDUCATION

MSc Artificial Intelligence and Robotics, University of Hertfordshire | 2025 – Sep 2026

BEng Electronics and Communication Engineering, East West Institute of Technology | 2014 – 2018

CERTIFICATIONS

- Certificate of Attendance, From Zero to Hero: Federated AI in Healthcare Systems, St John's College, Cambridge, Aug 2026
- ParsAI Workshop Certificate, Prompt Engineering, University of Hertfordshire (Dr. Abolfazl Zaraki), Apr 2026

Full right to work in the UK (Graduate Route visa from October 2026; no sponsorship required).